

# CSCI567\_Project\_StochasticResults

March 10, 2021

CSCI 567 - Machine Learning - Spring 2021

Project: Chekcpoint 1 - Data Preprocessing (03/10/21)

Team StochasticResults:

- Abel Salinas [8793999216] [abelsali@usc.edu]
- Angel Nieto [2211798052] [nietogar@usc.edu]
- Misael Morales [5832732058] [misaelmo@usc.edu]

---

The “Housing Price” dataset consists of 79 predictors for the house prices in Ames, Iowa. The training and testing set are already pre-split for us from the Kaggle version (<https://www.kaggle.com/c/house-prices-advanced-regression-techniques/overview>), and we will focus on different regression techniques to predict the housing SalePrice based on the given features.

This notebook focuses on the data pre-processing, wrangling, visualization, and statistical analysis. It is a crucial step in any machine-learning/data-analytics application to ensure proper data formatting in order to optimize the techniques implemented.

## 1 Table of Contents:

1. Load required packages
2. Data wrangling
3. Data visualization
4. Exploratory data analysis

---

## 2 1. Load required libraries

```
[1]: # Basic data management packages
import os
import numpy as np
import pandas as pd

# Visualization packages
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import pandas.plotting as pd_plot
%matplotlib inline
```

```
[2]: # Exploratory Data Analysis packages
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.mixture import GaussianMixture
from sklearn.neighbors import LocalOutlierFactor
```

```
[3]: # Regression and Modeling packages
import tensorflow as tf
import keras
from scipy import stats
from sklearn.metrics import mean_squared_error, r2_score

# Verify GPU compatibility
print("Tensorflow Version:", tf.__version__)
print("Tensorflow built with CUDA?", tf.test.is_built_with_cuda())
print(tf.config.list_physical_devices('CPU'))
print(tf.config.list_physical_devices('GPU'))
print("Num GPU Available:", len(tf.config.list_physical_devices('GPU')))
```

Using TensorFlow backend.

Tensorflow Version: 2.4.0

Tensorflow built with CUDA? True

[PhysicalDevice(name='/physical\_device:CPU:0', device\_type='CPU')]

[PhysicalDevice(name='/physical\_device:GPU:0', device\_type='GPU')]

Num GPU Available: 1

---

## 3 2. Data Wrangling

```
[4]: # Read CSV files for Train/Test datasets
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv')
```

After loading the data, we display the dataframe shapes for the training and testing sets. We notice that they are not the same shape, and therefore we will have to investigate the reason for this, and try to solve as best as possible. The original training set contains 1460 samples of 81 features, while the original testing set contains 1459 samples of 80 features. We can throw away the “Id” column, and obtain the true 79-feature datasets.

We also see that there are a lot of features that contain NaN’s, as well as a mixture of numeric and string types for the predictors. This will become problematic for some regression techniques, and therefore we must devise a technique to transform as best as possible all data into numeric.

```
[5]: print('Train Shape: {} | types: {} \nTest Shape: {} | types: {}'.
      ↪format(train_df.shape, pd.unique(train_df.dtypes),
                                                    test_df.
      ↪shape, pd.unique(test_df.dtypes)))
print('Set difference train-vs-test: {}'.format(set(train_df.columns).
      ↪difference(set(test_df.columns))))
```

```
Train Shape: (1460, 81) | types: [dtype('int64') dtype('O') dtype('float64')]
Test Shape: (1459, 80) | types: [dtype('int64') dtype('O') dtype('float64')]
Set difference train-vs-test: {'SalePrice'}
```

```
[6]: x_train = train_df.iloc[:,1:-1] #79 train features
y_train = train_df.iloc[:,-1] #SalePrice training target
x_test = test_df.iloc[:,1:] #79 test features
print('x_train {} | y_train {} \nx_test {}'.format(x_train.shape, y_train.
      ↪shape, x_test.shape))
```

```
x_train (1460, 79) | y_train (1460,)
x_test (1459, 79)
```

We preview the first 5 rows of the training set. This allows to explore the data type for some of the features provided in this set. From here, we realize that 79 features is a very large number of features to visually display each time, and so we will restrict future visualizations to the most important features, or the features of interests for the specific operations we are performing at the moment.

```
[7]: #preview the training data set
x_train.head()
```

```
[7]: MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0      60      RL      65.0      8450  Pave  NaN      Reg
1      20      RL      80.0      9600  Pave  NaN      Reg
2      60      RL      68.0     11250  Pave  NaN      IR1
3      70      RL      60.0      9550  Pave  NaN      IR1
4      60      RL      84.0     14260  Pave  NaN      IR1

LandContour Utilities LotConfig ... ScreenPorch PoolArea PoolQC Fence \
0      Lvl1    AllPub    Inside ...           0           0    NaN    NaN
1      Lvl1    AllPub      FR2 ...           0           0    NaN    NaN
2      Lvl1    AllPub    Inside ...           0           0    NaN    NaN
3      Lvl1    AllPub    Corner ...           0           0    NaN    NaN
4      Lvl1    AllPub      FR2 ...           0           0    NaN    NaN

MiscFeature MiscVal MoSold YrSold SaleType SaleCondition
0      NaN         0       2    2008        WD        Normal
1      NaN         0       5    2007        WD        Normal
2      NaN         0       9    2008        WD        Normal
3      NaN         0       2    2006        WD      Abnorml
4      NaN         0      12    2008        WD        Normal
```

[5 rows x 79 columns]

```
[8]: #preview the testing data set
x_test.head()
```

```
[8]: MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0      20      RH      80.0    11622  Pave   NaN    Reg
1      20      RL      81.0    14267  Pave   NaN    IR1
2      60      RL      74.0    13830  Pave   NaN    IR1
3      60      RL      78.0     9978  Pave   NaN    IR1
4     120      RL      43.0     5005  Pave   NaN    IR1

LandContour Utilities LotConfig ... ScreenPorch PoolArea PoolQC Fence \
0      Lvl1  AllPub   Inside ...      120      0   NaN  MnPrv
1      Lvl1  AllPub  Corner ...        0      0   NaN   NaN
2      Lvl1  AllPub   Inside ...        0      0   NaN  MnPrv
3      Lvl1  AllPub   Inside ...        0      0   NaN   NaN
4      HLS   AllPub   Inside ...     144      0   NaN   NaN

MiscFeature MiscVal  MoSold  YrSold  SaleType  SaleCondition
0      NaN        0        6   2010      WD      Normal
1     Gar2    12500        6   2010      WD      Normal
2      NaN        0        3   2010      WD      Normal
3      NaN        0        6   2010      WD      Normal
4      NaN        0        1   2010      WD      Normal
```

[5 rows x 79 columns]

We check to see how many of the features in the training set are categorical and how many are numerical. We see an almost even split, and therefore we are going to have to change the categorical predictors from strings into some sort of integer mapping.

```
[9]: numerical_feats = x_train.dtypes[x_train.dtypes != "object"].index
print("Number of Numerical features: ", len(numerical_feats))

categorical_feats = x_train.dtypes[x_train.dtypes == "object"].index
print("Number of Categorical features: ", len(categorical_feats))
```

```
Number of Numerical features: 36
Number of Categorical features: 43
```

We also count to see what percentage of the features are NaNs. For some variables, we observe a lot of missing data. This might indicate that it could be worth eliminating these features where the majority of the data is null.

```
[10]: total = x_train.isnull().sum().sort_values(ascending=False)
percent = (x_train.isnull().sum()/x_train.isnull().count()).
        sort_values(ascending=False)
```

```
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(10)
```

```
[10]:
```

	Total	Percent
PoolQC	1453	0.995205
MiscFeature	1406	0.963014
Alley	1369	0.937671
Fence	1179	0.807534
FireplaceQu	690	0.472603
LotFrontage	259	0.177397
GarageCond	81	0.055479
GarageType	81	0.055479
GarageYrBlt	81	0.055479
GarageFinish	81	0.055479

Let's now focus on fixing the training and testing features in order to remove NaN's, change strings and objects into numerics, and other important data preprocessing and wrangling operations. We will go one-by-one on the features that require some fixing or reinterpretation for future regression modeling.

```
[11]: # Non-numeric variables that require attention
non_num_vars = x_train.dtypes[x_train.dtypes=='object'].index
print(non_num_vars)
```

```
Index(['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',
      'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
      'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st',
      'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation',
      'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
      'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',
      'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual',
      'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature',
      'SaleType', 'SaleCondition'],
      dtype='object')
```

It is also important to know that some features have a slightly different definition than others in terms of their categorical values. For PoolQC, for example, NaN is not missing data but means no pool, likewise for Fence, FireplaceQu etc.

```
[12]: # Replace object/string values with categorial values from the
      ↪ "data_description" file from Kaggle
mappings = dict(
MSZoning_mapping      = {'nan':0, 'A':1, 'C (all)':2, 'FV':3, 'I':4, 'RH':5,
      ↪ 'RL':6, 'RP':7, 'RM':8},
Street_mapping        = {'nan':0, 'Grvl':1, 'Pave':2},
Alley_mapping          = {'nan':0, 'Grvl':1, 'Pave':2, 'NaN':0},
LotShape_mapping       = {'nan':0, 'Reg':1, 'IR1':2, 'IR2':3, 'IR3':4 },
LandContour_mapping   = {'nan':0, 'Lvl':1, 'Bnk':2, 'HLS':3, 'Low':4},
```

```

Utilities_mapping      = {'nan':0, 'AllPub':1, 'NoSewr':2, 'NoSeWa':3, 'ELO':4},
LotConfig_mapping      = {'nan':0, 'Inside':1, 'Corner':2, 'CulDSac':3, 'FR2':4,
↳ 'FR3':5},
LandSlope_mapping      = {'nan':0, 'Gtl':1, 'Mod':2, 'Sev':3},
Neighborhood_mapping   = {'nan':0, 'Blmngtn':1, 'Blueste':2, 'BrDale':3,
↳ 'BrkSide':4, 'ClearCr':5, 'CollgCr':6, 'Crawfor':7,
↳ 'Edwards':8, 'Gilbert':9, 'IDOTRR':10, 'MeadowV':11,
↳ 'Mitchel':12, 'Names':13, 'NoRidge':14,
↳ 'NPkVill':15, 'NridgHt':16, 'NWAmes':17, 'OldTown':18,
↳ 'SWISU':19, 'Sawyer':20, 'SawyerW':21,
↳ 'Somerst':22, 'StoneBr':23, 'Timber':24, 'Veenker':25},
Condition1_mapping     = {'nan':0, 'Artery':1, 'Feedr':2, 'Norm':3, 'RRNn':4,
↳ 'RRAn':5, 'PosN':6, 'PosA':7, 'RRNe':8, 'RAe':9},
Condition2_mapping     = {'nan':0, 'Artery':1, 'Feedr':2, 'Norm':3, 'RRNn':4,
↳ 'RRAn':5, 'PosN':6, 'PosA':7, 'RRNe':8, 'RAe':9},
BldgType_mapping       = {'nan':0, '1Fam':1, '2fmCon':2, 'Duplex':3, 'Twnhs':4,
↳ 'TwnhsE':4, 'TwnhsI':5},
HouseStyle_mapping     = {'nan':0, '1Story':1, '1.5Fin':2, '1.5Unf':3, '2Story':
↳ 4, '2.5Fin':5, '2.5Unf':6, 'SFoyer':7, 'SLvl':8},
RoofStyle_mapping      = {'nan':0, 'Flat':1, 'Gable':2, 'Gambrel':3, 'Hip':4,
↳ 'Mansard':5, 'Shed':6},
RoofMatl_mapping       = {'nan':0, 'ClyTile':1, 'CompShg':2, 'Membran':3,
↳ 'Metal':4,
↳ 'Roll':5, 'Tar&Grv':6, 'WdShake':7, 'WdShngl':8},
Exterior1st_mapping    = {'nan':0, 'AsbShng':1, 'AsphShn':2, 'BrkComm':3,
↳ 'BrkFace':4, 'CBlock':5, 'CemntBd':6,
↳ 'HdBoard':7, 'ImStucc':8, 'MetalSd':9, 'Other':10,
↳ 'Plywood':11, 'PreCast':12, 'Stone':13,
↳ 'Stucco':14, 'VinylSd':15, 'Wd Sdng':16, 'WdShing':17},
Exterior2nd_mapping    = {'nan':0, 'AsbShng':1, 'AsphShn':2, 'Brk Cmn':3,
↳ 'BrkFace':4, 'CBlock':5, 'CmentBd':6,
↳ 'HdBoard':7, 'ImStucc':8, 'MetalSd':9, 'Other':10,
↳ 'Plywood':11, 'PreCast':12, 'Stone':13,
↳ 'Stucco':14, 'VinylSd':15, 'Wd Shng':16, 'Wd Sdng':16,
↳ 'WdShing':17},
MasVnrType_mapping     = {'nan':0, 'BrkCmn':1, 'BrkFace':2, 'CBlock':3, 'None':
↳ 4, 'Stone':5},
ExterQual_mapping      = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5},
ExterCond_mapping      = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5},
Foundation_mapping     = {'nan':0, 'BrkTil':1, 'CBlock':2, 'PConc':3, 'Slab':4,
↳ 'Stone':5, 'Wood':6},
BsmtQual_mapping       = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5, 'NA':
↳ 0},
BsmtCond_mapping       = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5, 'NA':
↳ 0},
BsmtExposure_mapping   = {'nan':0, 'Gd':1, 'Av':2, 'Mn':3, 'No':4, 'NA':5},

```

```

BsmtFinType1_mapping = {'nan':0, 'GLQ':1, 'ALQ':2, 'BLQ':3, 'Rec':4, 'LwQ':5,
↳ 'Unf':6, 'NA':0},
BsmtFinType2_mapping = {'nan':0, 'GLQ':1, 'ALQ':2, 'BLQ':3, 'Rec':4, 'LwQ':5,
↳ 'Unf':6, 'NA':0},
Heating_mapping      = {'nan':0, 'Floor':1, 'GasA':2, 'GasW':3, 'Grav':4,
↳ 'OthW':5, 'Wall':6},
HeatingQC_mapping    = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5},
CentralAir_mapping   = {'nan':0, 'N':1, 'Y':2},
Electrical_mapping    = {'nan':0, 'SBrkr':1, 'FuseA':2, 'FuseF':3, 'FuseP':4,
↳ 'Mix':5},
KitchenQual_mapping  = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5},
Functional_mapping    = {'nan':0, 'Typ':1, 'Min1':2, 'Min2':3, 'Mod':4, 'Maj1':
↳ 5, 'Maj2':6, 'Sev':7, 'Sal':8},
FireplaceQu_mapping  = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5, 'NA':
↳ 0},
GarageType_mapping    = {'nan':0, '2Types':1, 'Attchd':2, 'Basment':3,
↳ 'BuiltIn':4, 'CarPort':5, 'Detchd':6, 'NA':0},
GarageFinish_mapping  = {'nan':0, 'Fin':1, 'RFn':2, 'Unf':3, 'NA':0},
GarageQual_mapping    = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5, 'NA':
↳ 0},
GarageCond_mapping    = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'Po':5, 'NA':
↳ 0},
PavedDrive_mapping    = {'nan':0, 'Y':1, 'P':2, 'N':3},
PoolQC_mapping        = {'nan':0, 'Ex':1, 'Gd':2, 'TA':3, 'Fa':4, 'NA':0},
Fence_mapping         = {'nan':0, 'GdPrv':1, 'MnPrv':2, 'GdWo':3, 'MnWw':4,
↳ 'NA':0},
MiscFeature_mapping    = {'nan':0, 'Elev':1, 'Gar2':2, 'Othr':3, 'Shed':4,
↳ 'TenC':5, 'NA':0},
SaleType_mapping      = {'nan':0, 'WD':1, 'CWD':2, 'VWD':3, 'New':4, 'COD':5,
↳ 'Con':6, 'ConLw':7, 'ConLI':8,
                        'ConLD':9, 'Oth':10},
SaleCondition_mapping = {'nan':0, 'Normal':1, 'Abnorml':2, 'AdjLand':3,
↳ 'Alloca':4, 'Family':5, 'Partial':6})

```

With the dictionaries for relabeling the object/string-types as integer classes, we can now recursively replace all of the unuseful features with more meaningful values. We also replace all NaN's with 0's, so that they don't have an impact on the future regression models.

```

[13]: x_train = x_train.replace({list(non_num_vars)[k] : list(mappings.values())[k]
                                for k in np.arange(len(list(non_num_vars)))}).
↳ fillna(0)
x_test  = x_test.replace({list(non_num_vars)[k] : list(mappings.values())[k]
                           for k in np.arange(len(list(non_num_vars)))}).
↳ fillna(0)

```

```

[14]: x_train.head(3)

```

```
[14]: MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  \
0          60          6          65.0    8450      2    0.0          1
1          20          6          80.0    9600      2    0.0          1
2          60          6          68.0   11250      2    0.0          2

      LandContour  Utilities  LotConfig  ...  ScreenPorch  PoolArea  PoolQC  \
0              1          1          1  ...           0          0    0.0
1              1          1          4  ...           0          0    0.0
2              1          1          1  ...           0          0    0.0

      Fence  MiscFeature  MiscVal  MoSold  YrSold  SaleType  SaleCondition
0     0.0          0.0          0        2    2008          1              1
1     0.0          0.0          0        5    2007          1              1
2     0.0          0.0          0        9    2008          1              1

[3 rows x 79 columns]
```

```
[15]: x_test.head(3)
```

```
[15]: MSSubClass  MSZoning  LotFrontage  LotArea  Street  Alley  LotShape  \
0          20          5.0          80.0   11622      2    0.0          1
1          20          6.0          81.0   14267      2    0.0          2
2          60          6.0          74.0   13830      2    0.0          2

      LandContour  Utilities  LotConfig  ...  ScreenPorch  PoolArea  PoolQC  \
0              1          1.0          1  ...        120          0    0.0
1              1          1.0          2  ...           0          0    0.0
2              1          1.0          1  ...           0          0    0.0

      Fence  MiscFeature  MiscVal  MoSold  YrSold  SaleType  SaleCondition
0     2.0          0.0          0        6    2010          1.0              1
1     0.0          2.0   12500        6    2010          1.0              1
2     2.0          0.0          0        3    2010          1.0              1

[3 rows x 79 columns]
```

We confirm that we are now rid of any NaN's and null values in the training and testing datasets.

```
[16]: print('Training set null values: %i' %x_train.isnull().sum().sum())
      print('Testing set null values: %i' %x_test.isnull().sum().sum())
```

```
Training set null values: 0
Testing set null values: 0
```

```
[17]: numerical_feats = x_train.dtypes[x_train.dtypes != "object"].index
      print("Number of Numerical features: ", len(numerical_feats))

      categorical_feats = x_train.dtypes[x_train.dtypes == "object"].index
```



```
print("Number of Categorical features: ", len(categorical_feats))
```

```
Number of Numerical features: 79  
Number of Categorical features: 0
```

---

## 4 3. Data Visualization

```
[18]: # Declare a function to plot a graphical correlation matrix  
def plot_corr(dataframe,size=8):  
    corr = dataframe.corr()  
    fig, ax = plt.subplots(figsize=(size, size))  
    im = ax.matshow(corr,vmin = -1.0, vmax = 1.0, cmap='seismic')  
    #plt.xticks(range(len(corr.columns)), corr.columns);  
    #plt.yticks(range(len(corr.columns)), corr.columns);  
    plt.xticks([]); plt.yticks([])  
    plt.colorbar(im, orientation = 'vertical', shrink=0.8)  
    plt.title('Correlation Matrix')
```

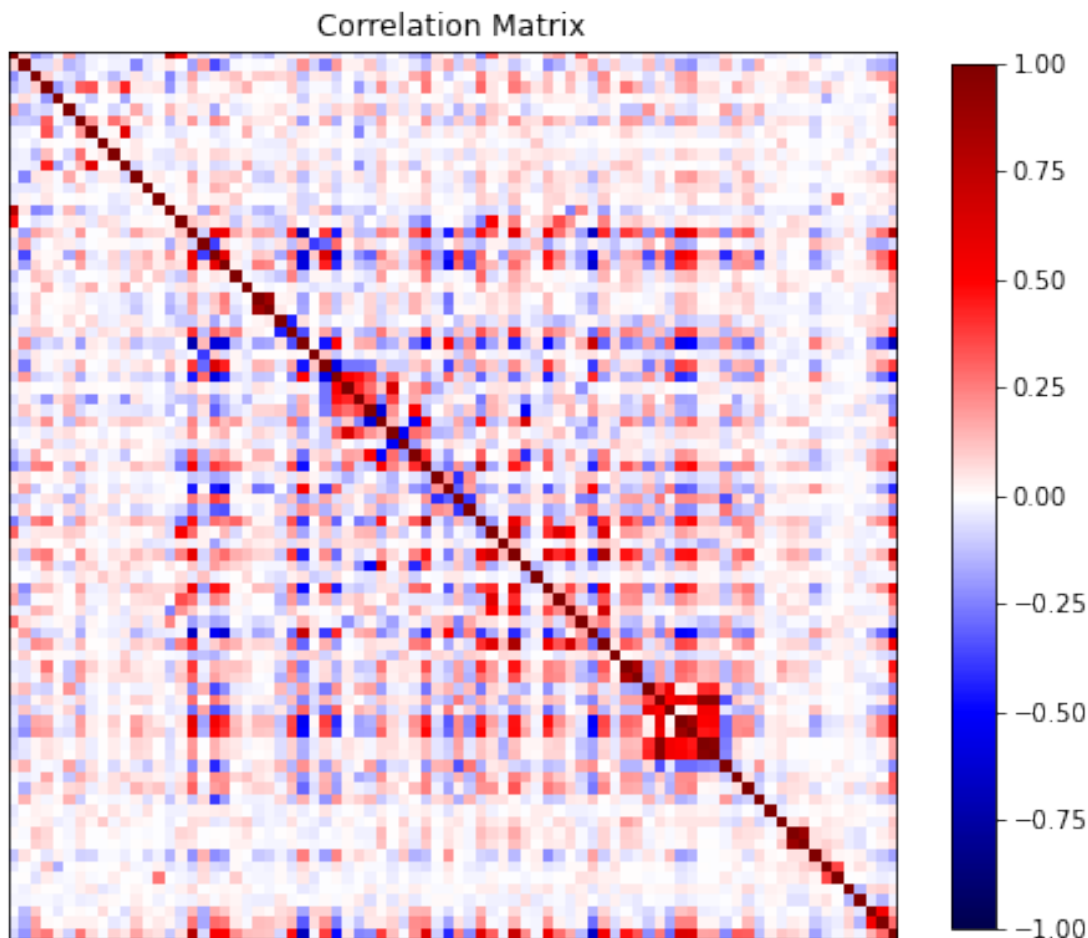
We start by visualizing the correlation between the 79 predictors and the target variable. This will tell us which predictors are highly correlated with each other, as well as which predictors have the highest correlation with the target SalePrice. We then select the last row (or column) of the matrix and select which predictors have the highest correlation to the target.

```
[19]: corr_mat = np.corrcoef(x_train.join(y_train),rowvar = False)  
corr_vec = corr_mat[:,-1] #correlation of features to target only  
print(corr_mat.round(2))
```

```
[[ 1.    0.08 -0.22 ... -0.    -0.04 -0.08]  
 [ 0.08  1.    -0.07 ... -0.13 -0.16 -0.21]  
 [-0.22 -0.07  1.    ...  0.11  0.18  0.21]  
 ...  
 [-0.    -0.13  0.11 ...  1.    0.54  0.15]  
 [-0.04 -0.16  0.18 ...  0.54  1.    0.29]  
 [-0.08 -0.21  0.21 ...  0.15  0.29  1.  ]]
```

The following correlation matrix/heatmap shows the effect of multicollinearity for some highly correlated variables that might be describing the same thing. For some cases, it might be important to note that we don't always need those pairs 2 or more variables that essentially describe the same thing - such as totalbsmtSF and 1stflrSF, or garageCars and garageArea.

```
[20]: plot_corr(x_train.join(y_train))
```



The user selects a level of correlation to cut-off; meaning that anything between this percentage will be considered as uncorrelated. For instance, a `corr_bound=0.7` means that anything between -0.7-to-0.7 is considered uncorrelated to SalePrice, while anything greater than 0.7 or less than -0.7 is considered correlated. We further explore these correlated features to the target variables in different ways.

```
[21]: #View features that have a +/- correlation greater than the user-specified
      ↪ correlation bound (corr_bound)
corr_bound = 0.6
corr_dat = (x_train.iloc[:, corr_vec<=-corr_bound]).join(x_train.iloc[:,
      ↪ corr_vec>corr_bound])
corr_dat
```

```
[21]:
```

	ExterQual	KitchenQual	OverallQual	TotalBsmtSF	1stFlrSF	GrLivArea	\
0	2	2	7	856	856	1710	
1	3	3	6	1262	1262	1262	
2	2	2	7	920	920	1786	

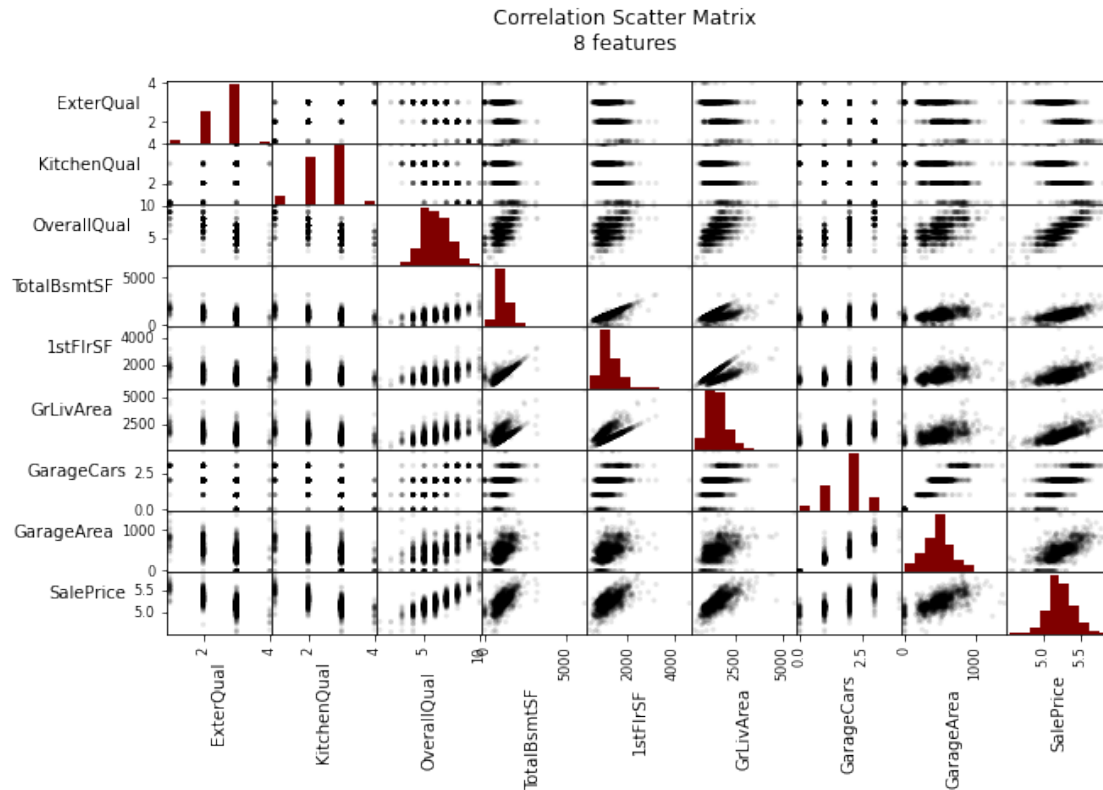
3	3	2	7	756	961	1717
4	2	2	8	1145	1145	2198
...	...	...	...	...	...	...
1455	3	3	6	953	953	1647
1456	3	3	6	1542	2073	2073
1457	1	2	7	1152	1188	2340
1458	3	2	5	1078	1078	1078
1459	2	3	5	1256	1256	1256

	GarageCars	GarageArea
0	2	548
1	2	460
2	2	608
3	3	642
4	3	836
...	...	...
1455	2	460
1456	2	500
1457	1	252
1458	1	240
1459	1	276

[1460 rows x 8 columns]

```
[22]: axes = pd_plot.scatter_matrix(corr_dat.join(np.log10(y_train)), alpha=0.1,
    figsize=(10, 6),
    color='black', hist_kwds={'color':['maroon']})

for ax in axes.flatten():
    ax.xaxis.label.set_rotation(90)
    ax.yaxis.label.set_rotation(0)
    ax.yaxis.label.set_ha('right')
plt.suptitle('Correlation Scatter Matrix\n{} features'.format(corr_dat.
    shape[-1]))
plt.show();
```



```
[23]: #full training set basic statistics
x_train.describe()
```

```
[23]:
```

	MSSubClass	MSZoning	LotFrontage	LotArea	Street	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	56.897260	6.126712	57.623288	10516.828082	1.995890	
std	42.300571	1.050330	34.664304	9981.264932	0.063996	
min	20.000000	2.000000	0.000000	1300.000000	1.000000	
25%	20.000000	6.000000	42.000000	7553.500000	2.000000	
50%	50.000000	6.000000	63.000000	9478.500000	2.000000	
75%	70.000000	6.000000	79.000000	11601.500000	2.000000	
max	190.000000	8.000000	313.000000	215245.000000	2.000000	

	Alley	LotShape	LandContour	Utilities	LotConfig	...	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	...	
mean	0.090411	1.408219	1.185616	1.001370	1.416438	...	
std	0.372151	0.582296	0.606509	0.052342	0.773448	...	
min	0.000000	1.000000	1.000000	1.000000	1.000000	...	
25%	0.000000	1.000000	1.000000	1.000000	1.000000	...	
50%	0.000000	1.000000	1.000000	1.000000	1.000000	...	
75%	0.000000	2.000000	1.000000	1.000000	2.000000	...	

max	2.000000	4.000000	4.000000	3.000000	5.000000	...
-----	----------	----------	----------	----------	----------	-----

	ScreenPorch	PoolArea	PoolQC	Fence	MiscFeature	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	15.060959	2.758904	0.010959	0.396575	0.144521	
std	55.757415	40.177307	0.177224	0.875914	0.742569	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	480.000000	738.000000	4.000000	4.000000	5.000000	

	MiscVal	MoSold	YrSold	SaleType	SaleCondition
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	43.489041	6.321918	2007.815753	1.490411	1.582192
std	496.123024	2.703626	1.328095	1.368616	1.475209
min	0.000000	1.000000	2006.000000	1.000000	1.000000
25%	0.000000	5.000000	2007.000000	1.000000	1.000000
50%	0.000000	6.000000	2008.000000	1.000000	1.000000
75%	0.000000	8.000000	2009.000000	1.000000	1.000000
max	15500.000000	12.000000	2010.000000	10.000000	6.000000

[8 rows x 79 columns]

```
[24]: #main features (based on correlation coefficient) basic statistics
corr_dat.describe()
```

```
[24]:
```

	ExterQual	KitchenQual	OverallQual	TotalBsmtSF	1stFlrSF	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	
mean	2.60411	2.488356	6.099315	1057.429452	1162.626712	
std	0.57428	0.663760	1.382997	438.705324	386.587738	
min	1.00000	1.000000	1.000000	0.000000	334.000000	
25%	2.00000	2.000000	5.000000	795.750000	882.000000	
50%	3.00000	3.000000	6.000000	991.500000	1087.000000	
75%	3.00000	3.000000	7.000000	1298.250000	1391.250000	
max	4.00000	4.000000	10.000000	6110.000000	4692.000000	

	GrLivArea	GarageCars	GarageArea
count	1460.000000	1460.000000	1460.000000
mean	1515.463699	1.767123	472.980137
std	525.480383	0.747315	213.804841
min	334.000000	0.000000	0.000000
25%	1129.500000	1.000000	334.500000
50%	1464.000000	2.000000	480.000000
75%	1776.750000	2.000000	576.000000
max	5642.000000	4.000000	1418.000000

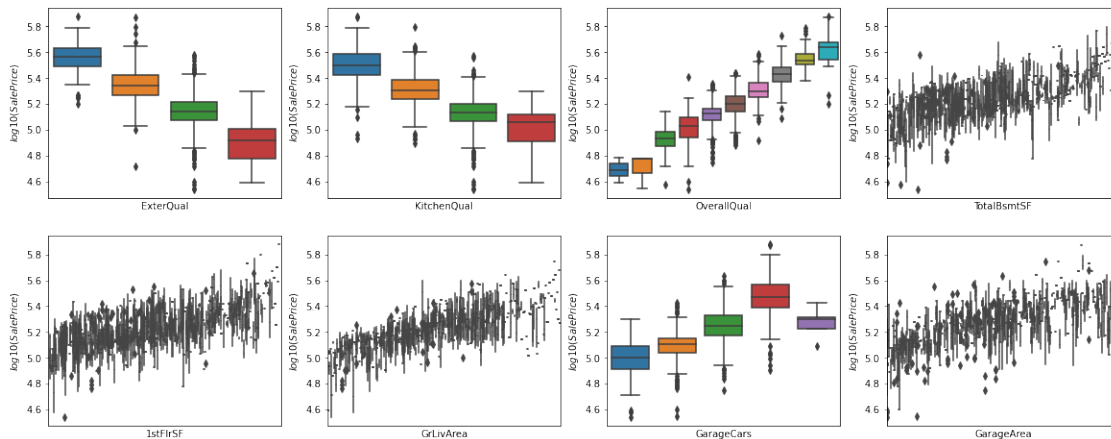
With these N correlated features, we now explore how the strong is the correlation with the SalePrice

target variable. We observe the numerical value, as well as the boxplot behavior of these highly-correlated features.

```
[25]: corr_dat.join(y_train).corr()['SalePrice']
```

```
[25]: ExterQual    -0.682639
      KitchenQual -0.659600
      OverallQual  0.790982
      TotalBsmtSF  0.613581
      1stFlrSF     0.605852
      GrLivArea     0.708624
      GarageCars    0.640409
      GarageArea    0.623431
      SalePrice     1.000000
      Name: SalePrice, dtype: float64
```

```
[26]: plt.figure(figsize=(20,corr_dat.shape[-1]))
      for k in np.arange(corr_dat.shape[-1]):
          plt.subplot(corr_dat.shape[-1]//4,4,k+1)
          sns.boxplot(x=corr_dat.iloc[:,k], y=np.log10(y_train))
          plt.xticks([]); plt.ylabel('$log10(SalePrice)$')
      plt.show();
```



We also visualize the target variable, SalePrice. We notice that these values are log-normally distributed. There are a very small amount of extremely expensive houses, while the vast majority are closer to the average. Therefore, taking the logarithm of this will make them Normally distributed. This will help a lot for the regression techniques, and therefore we will be working with the natural log of SalePrice from now on.

```
[27]: target = np.log10(y_train)
      print('SalePrice: Mean {:.3e} | Std. Dev {:.3e}'.format(y_train.mean(),
      ↪ y_train.std()))
```

```
print('Log10(SalePrice): Mean {:.3f}      | Std. Dev {:.3f}'.format(target.
↪mean(), target.std()))
```

```
SalePrice: Mean 1.809e+05 | Std. Dev 7.944e+04
Log10(SalePrice): Mean 5.222      | Std. Dev 0.173
```

```
[28]: print('      SalePrice: Skewness={:.3f} | Kurtosis={:.3f}'.format(y_train.
↪skew(), y_train.kurt()))
print('Log10(SalePrice): Skewness={:.3f} | Kurtosis={:.3f}'.format(target.
↪skew(), target.kurt()))
```


```
SalePrice: Skewness=1.883 | Kurtosis=6.536
Log10(SalePrice): Skewness=0.121 | Kurtosis=0.810
```

```
[29]: plt.figure(figsize=(20,5))

ax1 = plt.subplot(121)
plt.title('$SalePrice$ Distribution')
ax1.text(5E5,6E-6, '={:.2e}\nσ={:.2e}'.format(y_train.mean(), y_train.std()),␣
↪fontSize=12)
plt.hist(y_train, density=True, bins=20, alpha=0.7, histtype='bar', ec='black')
sns.kdeplot(y_train, linewidth=3, c='r')

ax2 = plt.subplot(122)
plt.title('$Log(SalePrice)$ Distribution')
ax2.text(5.5, 2, '={:.2f}\nσ={:.2f}'.format(target.mean(), target.std()),␣
↪fontSize=12)
plt.hist(target, density=True, bins=20, alpha=0.7, histtype='bar', ec='black')
sns.kdeplot(target, linewidth=3, c='r')

plt.show();
```



On a slightly different note, we will plot the relationship between two variables at a time and the target  $\log_{10}(\text{SalePrice})$ , using Seaborn, and try to fit a linear regressor onto the data. This will test for linearity and for possible colinearity, and will also compute the  $r^2$  coefficient and  $p$  value for the regression estimate.

```

[30]: nr_rows, nr_cols = 12, 4
       colors = sns.color_palette()

       fig, axs = plt.subplots(nr_rows, nr_cols, figsize=(nr_cols*3.5,nr_rows*3))

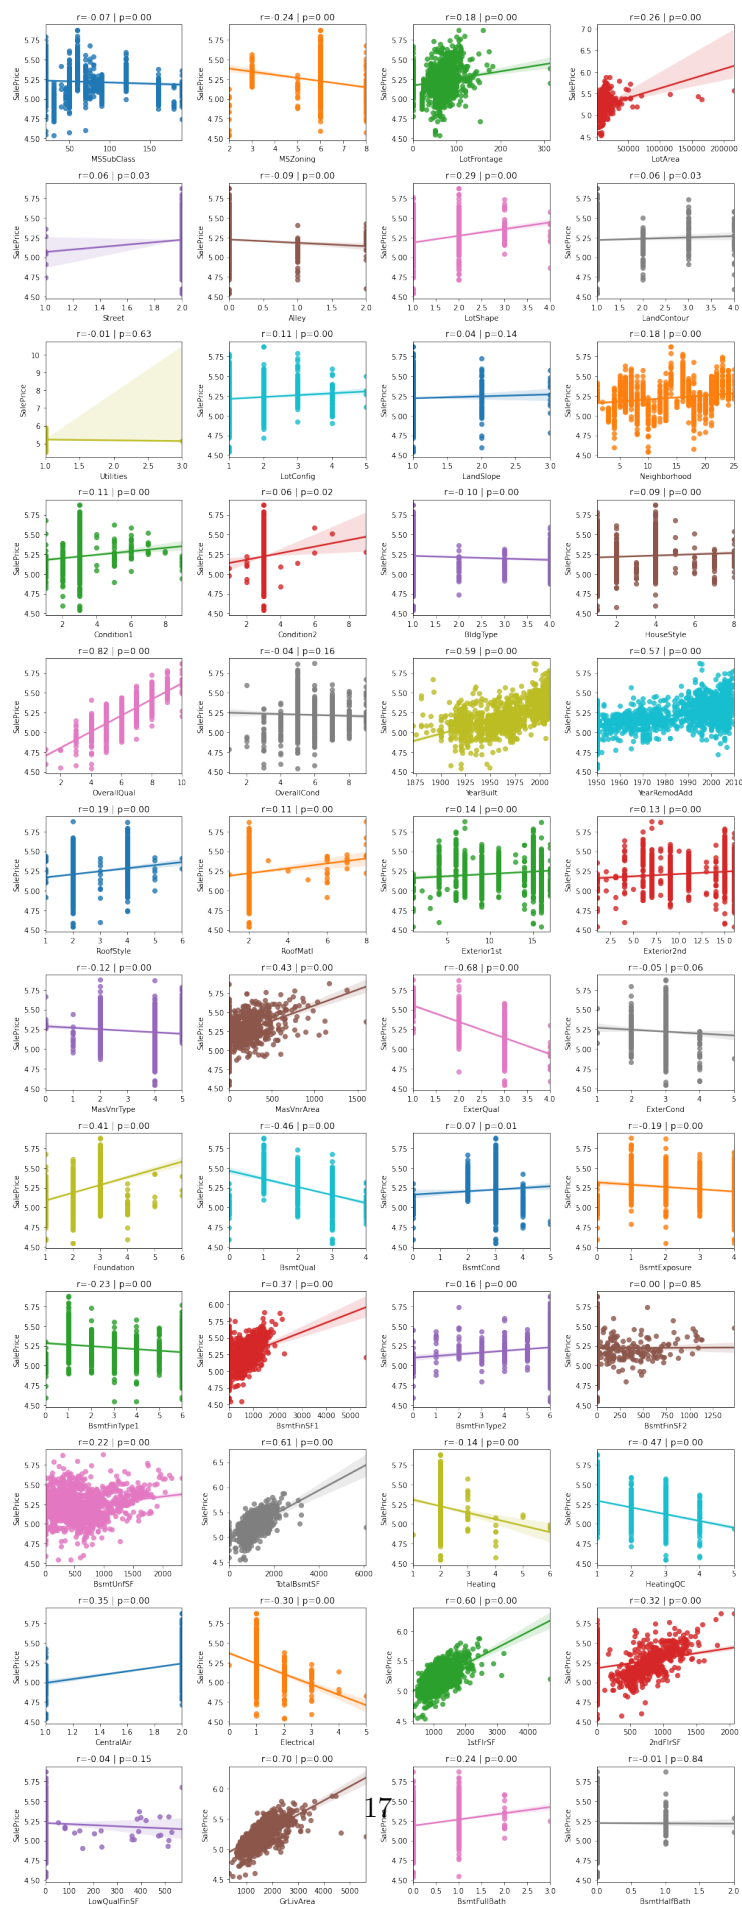
       li_num_feats = list(numerical_feats)
       li_not_plot = ['Id', 'SalePrice', 'SalePrice_Log']
       li_plot_num_feats = [c for c in list(numerical_feats) if c not in li_not_plot]

       for r in range(0,nr_rows):
           for c in range(0,nr_cols):
               i = r*nr_cols+c
               if i < len(li_plot_num_feats):
                   sns.regplot(x=x_train[li_plot_num_feats[i]], y=target,
                               ax = axs[r][c], color=colors[i % 10])
                   stp = stats.pearsonr(x_train[li_plot_num_feats[i]], target)
                   str_title = "r="+"{0:.2f}".format(stp[0])+" | "+"p="+"{0:.2f}".
                               ↪format(stp[1])
                   axs[r][c].set_title(str_title,fontsize=12)

       plt.tight_layout()
       plt.show()

```





---

## 5 4. Exploratory Data Analysis

After understanding the data statistics and relationships, we will now try to do some basic data analytics on the preprocessed data.

```
[31]: def plot_corr_matrix(df, nr_c, targ) :  
    corr_abs = df.corr().abs()  
    cols = corr_abs.nlargest(nr_c, targ)[targ].index  
    cm = np.corrcoef(df[cols].values.T)  
    plt.figure(figsize=(nr_c, nr_c))  
    plt.title('Correlation Matrix')  
    sns.set(font_scale=1.25)  
    sns.heatmap(cm, linewidths=1.5, annot=True, square=True,  
                fmt='.2f', annot_kws={'size': 10}, cbar_kws={'shrink':0.8},  
                yticklabels=cols.values, xticklabels=cols.values)  
    plt.show();
```

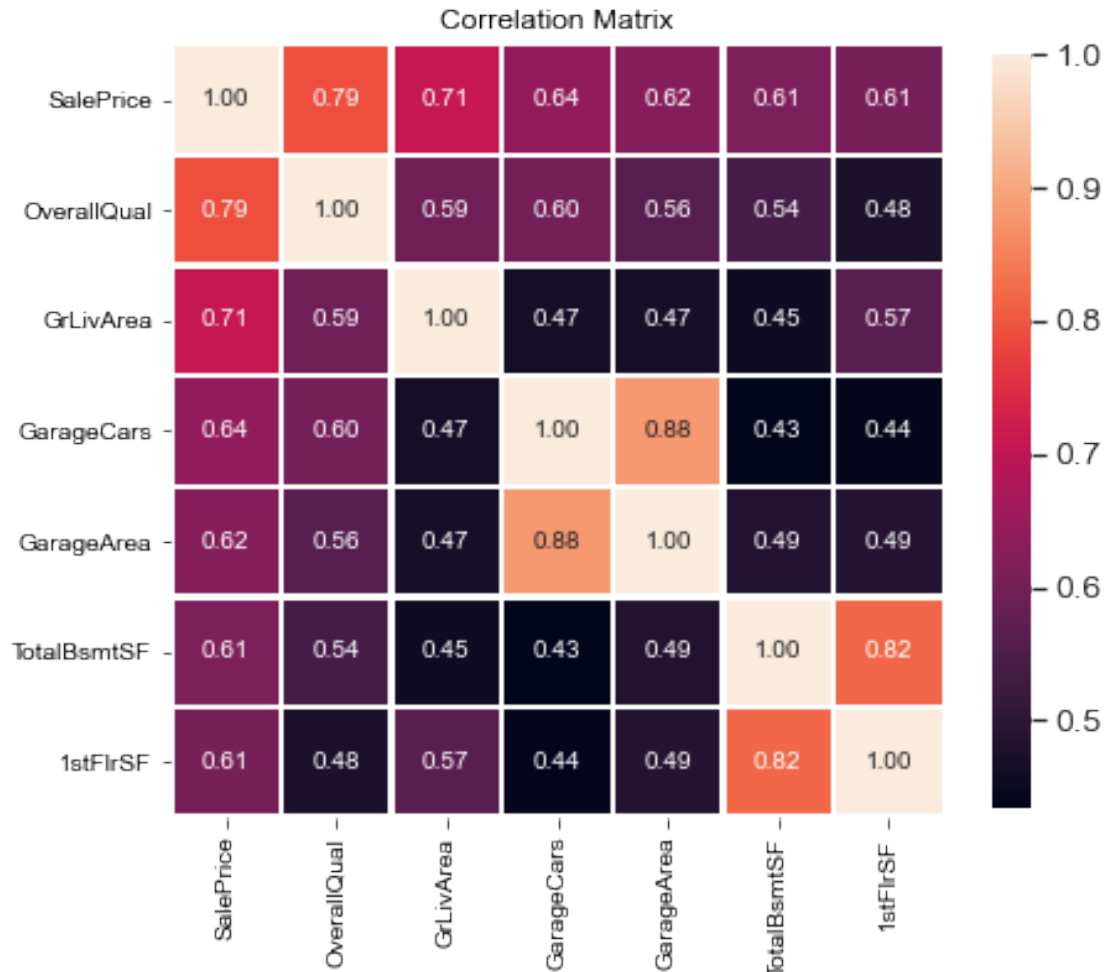
```
[32]: print('Significant Correlation Bound: ', corr_bound)  
    print('Shape of training data above correlation bound: ', corr_dat.shape)
```

Significant Correlation Bound: 0.6

Shape of training data above correlation bound: (1460, 8)

```
[33]: target_name = 'SalePrice'  
  
corr_abs = train_df.corr().abs()  
ser_corr = corr_abs.nlargest(len(numerical_feats), target_name)[target_name]  
  
cols_abv_corr_limit = list(ser_corr[ser_corr.values >= corr_bound].index)  
cols_bel_corr_limit = list(ser_corr[ser_corr.values < corr_bound].index)  
  
nr_feats = len(cols_abv_corr_limit)
```

```
[34]: plot_corr_matrix(train_df, nr_feats, target_name)
```



From the covariance matrix, we see that for some features like ‘OverallQual’ there is a strong linear correlation (0.8 approximately) to the target. For other features like ‘MSSubClass’ or ‘MiscFeature’ the correlation is very weak.

Scale the correlation-based data so that all features values are between 0 and 1. This will help greatly for the regression modeling and the visualization and interpretation of the models and the data.

```
[35]: scaler = StandardScaler()
      x_train_s = scaler.fit_transform(corr_dat)
```

Perform PCA to encode the correlation-based data into orthogonal singular vectors that will best describe the variability of the data. This will help visualize and try and classify clusters of data based on these new principal directions.

```
[36]: n_components = 4
      pca = PCA(n_components=n_components)
```

```
pca.fit(x_train_s)
```

```
#variance explained by each PC
```

```
print('PC Variance Explained: {}'.format(pca.explained_variance_ratio_))
```

```
print('Cumulative Variance Explained: {}'.format(np.sum(pca.  
→explained_variance_ratio_).round(4)))
```

PC Variance Explained: [0.59394484 0.12746352 0.10382664 0.07290565]

Cumulative Variance Explained: 0.8981

We see that with 4 PC's, for instance, we can explain approximately 90% of the variance in the scaled, clean data set.

[37]: *#plot n\_PC-vs-varExplained*

```
plt.figure()
```

```
plt.plot(range(len(np.cumsum(pca.explained_variance_ratio_))), np.cumsum(pca.  
→explained_variance_ratio_),
```

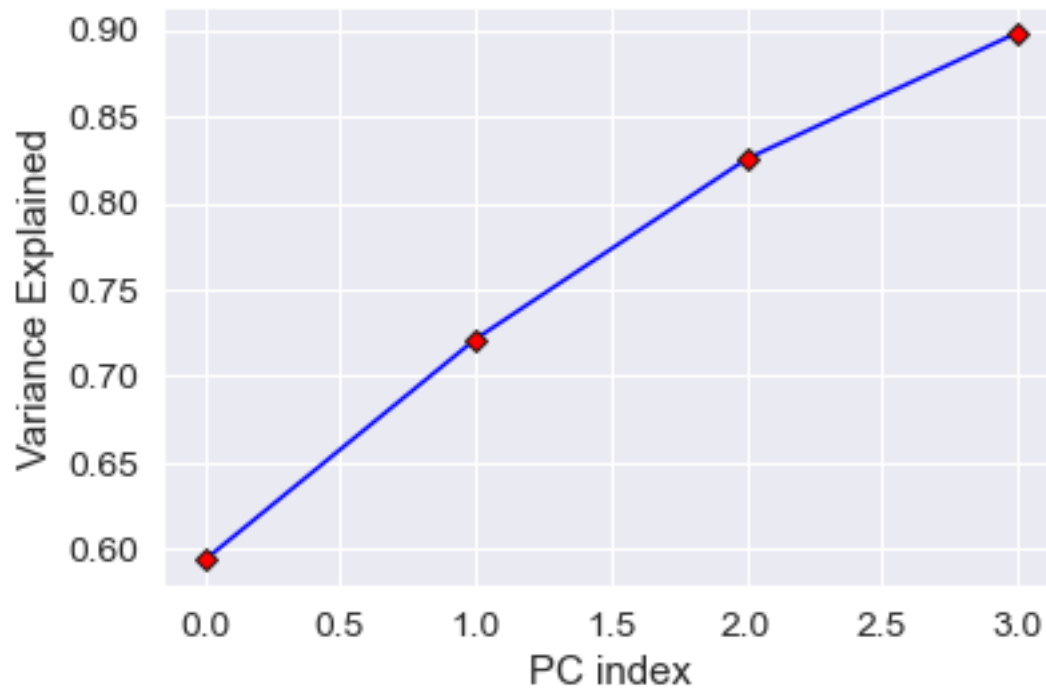
```
      '-bD', c='blue', mfc='red', mec='k')
```

```
plt.xlabel('PC index')
```

```
plt.ylabel('Variance Explained')
```

```
plt.grid('on')
```


```
plt.show();
```



We can see that PC1 and PC3 help us find a slight classification in PC-space for the categorical variables 'ExterQual', 'KitchenQual', 'OverallQual', and 'GarageCars'; especially the last one.

```
[38]: x_trans = pca.transform(x_train_s)

plt.figure(figsize=(20,6))
plt.subplot(221)
plt.scatter(x_trans[:,0], x_trans[:,2], c=corr_dat['ExterQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC ExterQual'); plt.xlabel('PC1'); plt.ylabel('PC3'); plt.xticks([]);
            ↪ plt.yticks([])
plt.subplot(222)
plt.scatter(x_trans[:,0], x_trans[:,2], c=corr_dat['KitchenQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC KitchenQual'); plt.xlabel('PC1'); plt.ylabel('PC3'); plt.
            ↪xticks([]); plt.yticks([])
plt.subplot(223)
plt.scatter(x_trans[:,0], x_trans[:,2], c=corr_dat['OverallQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC OverallQual'); plt.xlabel('PC1'); plt.ylabel('PC3'); plt.
            ↪xticks([]); plt.yticks([])
plt.subplot(224)
plt.scatter(x_trans[:,0], x_trans[:,2], c=corr_dat['GarageCars'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC GarageCars'); plt.xlabel('PC1'); plt.ylabel('PC3'); plt.
            ↪xticks([]); plt.yticks([])
plt.show();
```



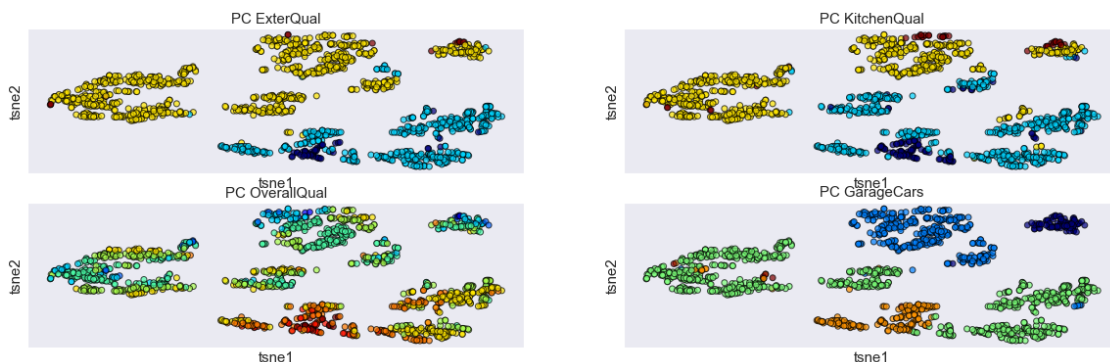
On the other hand, we can perform t-distributed Stochastic Neighbor Embedding (TSNE), which converts similarities between data points to joint probabilities and tries to minimize the Kullback-Leibler divergence between the joint probabilities of the low-dimensional embedding and the high-dimensional data. Here, we can see a good separation for ‘ExterQual’ and ‘GarageCars’ features, proving useful for predicting sale price using these variables.

```
[39]: tsne = TSNE(n_components = 2).fit_transform(x_train_s)
```

```

plt.figure(figsize=(20,6))
plt.subplot(221)
plt.scatter(tsne[:,0], tsne[:,1], c=corr_dat['ExterQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC ExterQual'); plt.xlabel('tsne1'); plt.ylabel('tsne2'); plt.
            ↪xticks([]); plt.yticks([])
plt.subplot(222)
plt.scatter(tsne[:,0], tsne[:,1], c=corr_dat['KitchenQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC KitchenQual'); plt.xlabel('tsne1'); plt.ylabel('tsne2'); plt.
            ↪xticks([]); plt.yticks([])
plt.subplot(223)
plt.scatter(tsne[:,0], tsne[:,1], c=corr_dat['OverallQual'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC OverallQual'); plt.xlabel('tsne1'); plt.ylabel('tsne2'); plt.
            ↪xticks([]); plt.yticks([])
plt.subplot(224)
plt.scatter(tsne[:,0], tsne[:,1], c=corr_dat['GarageCars'], cmap='jet',
            ↪edgecolors="black", alpha=0.7)
plt.title('PC GarageCars'); plt.xlabel('tsne1'); plt.ylabel('tsne2'); plt.
            ↪xticks([]); plt.yticks([])
plt.show();

```



Next, we use K-Means Clustering and Hierarchical Clustering to try and find some relationships in the data. This algorithm is unsupervised, and therefore tries to cluster the data based on similar predictor values toward the target. This will allow us to observe how we can start classifying and relating the data to the SalePrice. Using 3 clusters we might be able to visualize high/medium/low price houses based on these different highly-correlated predictor features.

```

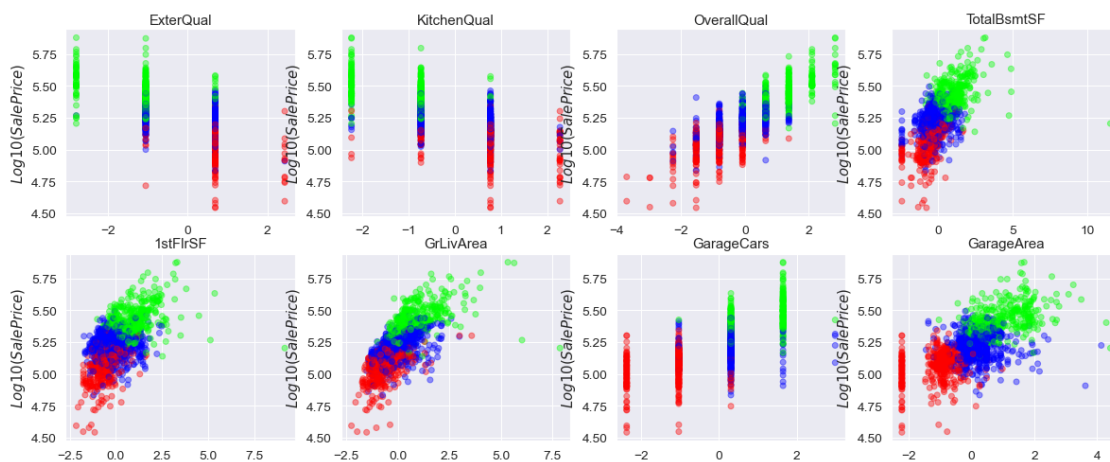
[40]: n_clusters = 3
kmeans = KMeans(n_clusters = n_clusters).fit(X=x_train_s, y=target).labels_
hclust = AgglomerativeClustering(n_clusters = n_clusters).fit(X=x_train_s,
            ↪y=target).labels_

```

```
gmm = GaussianMixture(n_components=n_clusters).fit_predict(X=x_train_s,
↳ y=target)
```

```
[41]: #select the clustering method to visualize
clust_method = kmeans

plt.figure(figsize=(20,8))
for k in range(len(corr_dat.columns)):
    plt.subplot(x_train_s.shape[-1]//4,4,k+1)
    plt.title('{}'.format(corr_dat.columns[k]))
    plt.scatter(x_train_s[:,k], target, c=clust_method, cmap='brg', alpha=0.4)
    plt.ylabel('$Log10(SalePrice)$')
plt.show();
```



We now perform Local Outlier Factor analysis to try and see if we can find any major outliers in the highly-correlated subset of the data.

```
[42]: lof = LocalOutlierFactor(n_neighbors=5)
lof.fit_predict(X=x_train_s, y=target)
lof_factors = -lof.negative_outlier_factor_
print('Local Outlier Factors: {}'.format(lof_factors)) #outliers are further_
↳ away from 1

# let outliers be those that are more than this upper bound value
lof_outliers_bound = 5
```

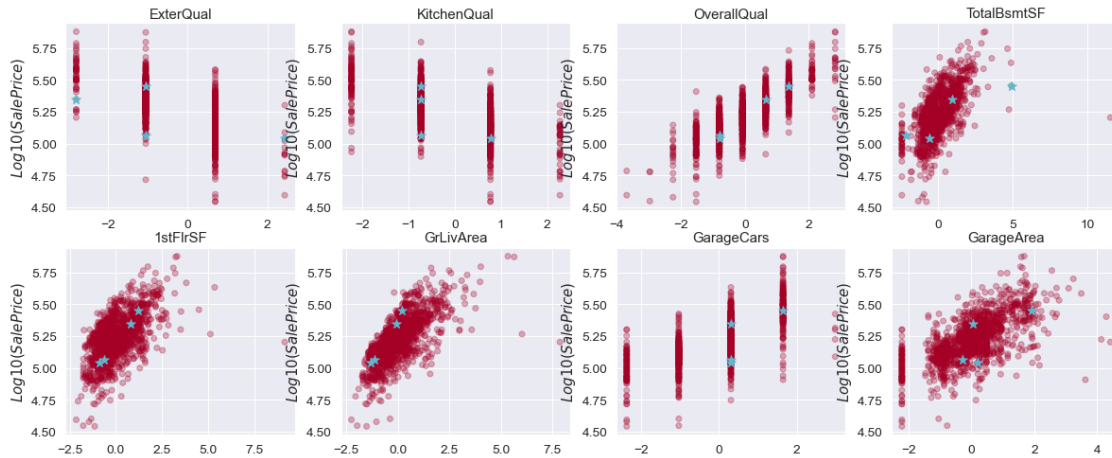
```
Local Outlier Factors: [0.9848064  0.97417749 1.10287748 ... 2.64963468
0.95880616 1.08014456]
```

```
[43]: plt.figure(figsize=(20,8))
for k in range(len(corr_dat.columns)):
    plt.subplot(x_train_s.shape[-1]//4,4,k+1)
```

```

plt.title('{}'.format(corr_dat.columns[k]))
plt.scatter(x_train_s[:,k], target, c=lof_factors>lof_outliers_bound,
cmap='RdYlGn', alpha=0.3)
plt.scatter(x_train_s[:,k][lof_factors>lof_outliers_bound],
target[lof_factors>lof_outliers_bound],
c='c', alpha=1, marker="*", s=100)
plt.ylabel('$Log10(SalePrice)$')
plt.show();

```



### Concluding Remarks:

- Remarks
  - The data has a different types of predictor features, including categorical and numerical. We mapped all categorical features into classes, which can then be used for supervised learning techniques.
  - The target variables, SalePrice, is log-normally distributed and thus any regression algorithms should be predicting the  $\log_{10}(\text{SalePrice})$  pseudofeature instead.
  - Data preprocessing and analysis included statistical and visual analysis; as well as unsupervised learning, also known as clustering analysis.
  - Some outliers are still present in the data, though they are only but a small percentage of the instances in the highly-correlated features selected.
- Conclusion
  - The Ames Housing data set is a complex yet versatile data set from machine learning projects.
  - Preprocessing and data wrangling is crucial for exploiting the full potential of all features.
  - The data is now ready to be used in different regression models to predict SalePrice from the features.
- Part 2: Regression We will attempt to use 6 different regression algorithms on the preprocessed data set. These potentially include linear regression, regularized regression, boosting,



SVM, random forest, and neural networks. We will then compare the performance of the different algorithms in predicting the (logarithm) SalePrice using the (subset) 79 (preprocessed) predictors provided. The mean-squared error metric will be used for optimization, and we hope to produce very accurate and interesting predictive models for the Ames Housing dataset.

---

## 6 End of Notebook